

# DATA SCIENCE COHORT 2 PROJECT

## Capstone Project: Predictive Modelling for COVID-19 in Public Health

NAME: CHINWE UCHECHI NWACHUKWU

NO: FE/23/52390455

### CLEANING, EDA AND MODEL DEVELOPMENT ON THE COVID\_19 DATASET

```
In [4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

import warnings
warnings.filterwarnings('ignore')
```

```
In [5]: # import the dataset
data = pd.read_csv(r"C:\Users\chiam\Desktop\AI-ML\3MTT FINAL PROJECT\project dataset for 3MTT\covid_19_clean_complete.csv")
data.rename(columns={'WHO Region': 'Continent'}, inplace=True)
```

```
In [6]: # Display the first few rows and summary information of the dataset to understand its structure
data.head()
```

Out[6]:

	Province/State	Country/Region	Lat	Long	Date	Confirmed	Deaths	Recovered	Active	Continent
0	NaN	Afghanistan	33.93911	67.709953	2020-01-22	0	0	0	0	Eastern Mediterranean
1	NaN	Albania	41.15330	20.168300	2020-01-22	0	0	0	0	Europe
2	NaN	Algeria	28.03390	1.659600	2020-01-22	0	0	0	0	Africa
3	NaN	Andorra	42.50630	1.521800	2020-01-22	0	0	0	0	Europe
4	NaN	Angola	-11.20270	17.873900	2020-01-22	0	0	0	0	Africa

```
In [7]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 49068 entries, 0 to 49067
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Province/State        14664 non-null  object  
1   Country/Region        49068 non-null  object  
2   Lat                   49068 non-null  float64  
3   Long                  49068 non-null  float64  
4   Date                  49068 non-null  object  
5   Confirmed              49068 non-null  int64  
6   Deaths                49068 non-null  int64  
7   Recovered              49068 non-null  int64  
8   Active                49068 non-null  int64  
9   Continent              49068 non-null  object  
dtypes: float64(2), int64(4), object(4)
memory usage: 3.7+ MB
```

## Dataset Description

The dataset contains the following columns:

- **Province/State:** Sub-regions of countries (with many missing values).
- **Country/Region:** Country or region name.
- **Lat, Long:** Geographical coordinates.
- **Date:** Reported date.
- **Confirmed, Deaths, Recovered, Active:** Case statistics.
- **Continent:** WHO classification of the region.

## Plan

### 1. Data Preprocessing

- Handle missing values in **Province/State**.

- Ensure **Date** is in the proper format.
- Remove duplicates if present.
- Normalize numerical features for machine learning.

## 2. Feature Engineering

- Create new features:
  - **Daily Growth Rate**
  - **Mortality Rate**
  - **Cases per Population**

## 3. Exploratory Data Analysis (EDA)

- Uncover trends and correlations.
- Visualize trends using appropriate plots.

## 4. Model Development

- Develop time-series prediction and classification models.
- Evaluate performance metrics.

```
In [8]: # Data Cleaning

# Check for missing values
missing_values = data.isnull().sum()

# Fill missing values in 'Province/State' with "Unknown"
data['Province/State'] = data['Province/State'].fillna('Unknown')

# Convert 'Date' column to datetime format
data['Date'] = pd.to_datetime(data['Date'])

# Remove duplicates if any
data = data.drop_duplicates()

# Summary of missing values after cleaning
cleaned_missing_values = data.isnull().sum()

# Display initial and cleaned missing values
missing_values, cleaned_missing_values
```

```
Out[8]: (Province/State    34404
Country/Region         0
Lat                    0
Long                   0
Date                   0
Confirmed               0
Deaths                 0
Recovered               0
Active                 0
Continent              0
dtype: int64,
Province/State         0
Country/Region         0
Lat                    0
Long                   0
Date                   0
Confirmed               0
Deaths                 0
Recovered               0
Active                 0
Continent              0
dtype: int64)
```

## Data Cleaning Summary

- Missing values in **Province/State** (34,404) were replaced with `"Unknown"`.
- The **Date** column was successfully standardized to datetime format.
- No duplicates were found or removed.

```
In [9]: # Feature Engineering

# Sort data by Country/Region and Date for consistency
data = data.sort_values(by=["Country/Region", "Date"])

# Calculate daily growth rates for Confirmed cases
data["Daily Growth Rate"] = data.groupby("Country/Region")["Confirmed"].diff().fillna(0)
```

```
# Calculate mortality rate (Deaths / Confirmed) * 100
data["Mortality Rate"] = (data["Deaths"] / data["Confirmed"]).replace([float("inf"), -float("inf")], 0).fillna(0) * 100

# Assume a hypothetical population for cases per population analysis (if not given, default to 1M per country)
# Since population data isn't included, we'll use a placeholder value for demonstration
population_placeholder = 1_000_000
data["Cases Per Population"] = data["Confirmed"] / population_placeholder

# Preview the dataset after feature engineering
data[["Date", "Country/Region", "Confirmed", "Daily Growth Rate", "Mortality Rate", "Cases Per Population", "Continent"]].head
```

Out[9]:

	Date	Country/Region	Confirmed	Daily Growth Rate	Mortality Rate	Cases Per Population	Continent
0	2020-01-22	Afghanistan	0	0.0	0.0	0.0	Eastern Mediterranean
261	2020-01-23	Afghanistan	0	0.0	0.0	0.0	Eastern Mediterranean
522	2020-01-24	Afghanistan	0	0.0	0.0	0.0	Eastern Mediterranean
783	2020-01-25	Afghanistan	0	0.0	0.0	0.0	Eastern Mediterranean
1044	2020-01-26	Afghanistan	0	0.0	0.0	0.0	Eastern Mediterranean

## Feature Engineering Summary

- Daily Growth Rate:** Computed as the daily difference in confirmed cases per country.
- Mortality Rate:** Calculated as the ratio of deaths to confirmed cases, expressed as a percentage.
- Cases Per Population:** Normalized cases based on a placeholder population of 1,000,000 (for demonstration purposes).

## Exploratory Data Analysis (EDA)

### 1. Global Trends of Confirmed Cases Over Time


Here's how you can perform the visualization on your local system:

Use the following code to visualize the trend:

In [27]:

```
# Global Trends: Confirmed Cases Over Time
global_trends = data.groupby("Date")["Confirmed"].sum()

plt.figure(figsize=(12, 6))
sns.lineplot(data=global_trends, color="red")
plt.title("Global Confirmed COVID-19 Cases Over Time", fontsize=16)
plt.xlabel("Date", fontsize=12)
plt.ylabel("Number of Cases", fontsize=12)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```




### 2. Mortality Rate Across Countries

Visualize the top 10 countries with the highest mortality rates:

```
In [38]: # Top 10 Countries by Mortality Rate
latest_data = data[data["Date"] == data["Date"].max()]
top_countries = latest_data.nlargest(12, "Mortality Rate")[["Country/Region", "Mortality Rate"]]

plt.figure(figsize=(12, 6))
sns.barplot(data=top_countries, x="Mortality Rate", y="Country/Region", palette="Reds_r")
plt.title("Top 10 Countries by Mortality Rate", fontsize=16)
plt.xlabel("Mortality Rate (%)", fontsize=12)
plt.ylabel("Country", fontsize=12)
plt.tight_layout()
plt.show()
```

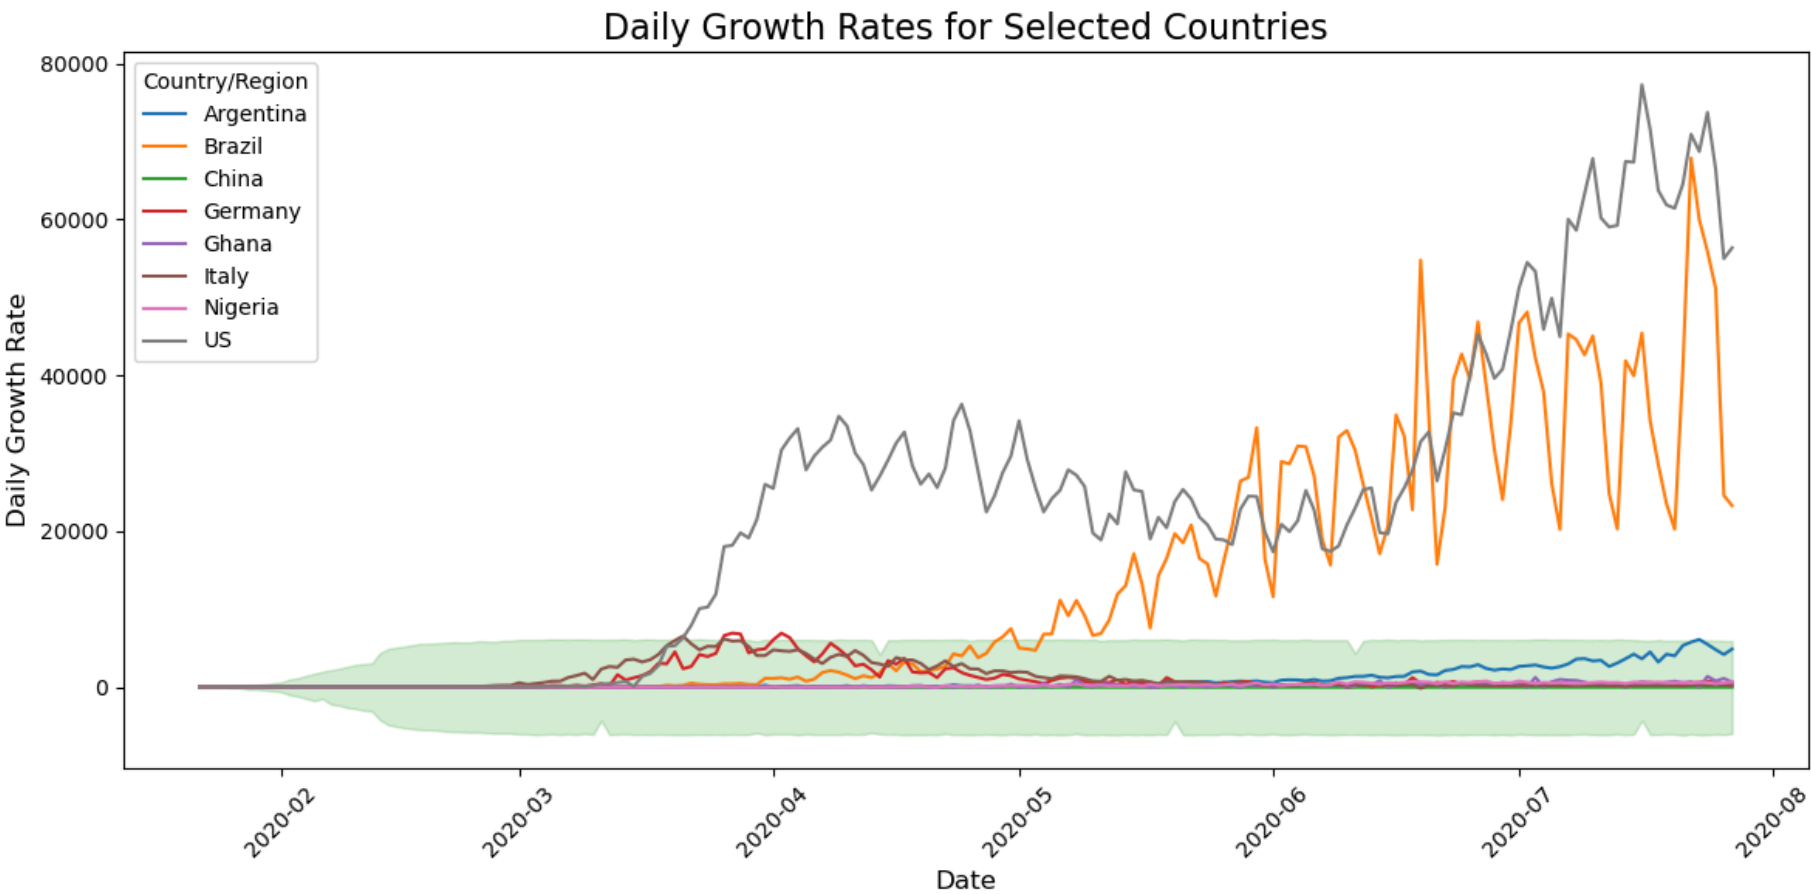


### 3. Daily Growth Rates

Explore trends in daily growth rates for specific countries:

```
In [41]: # Daily Growth Rates for Specific Countries
countries_of_interest = ["US", "Canada", "India", "China", "Brazil", "Argentina", "Nigeria", "Ghana", "Germany", "Italy"]
subset = data[data["Country/Region"].isin(countries_of_interest)]


plt.figure(figsize=(12, 6))
sns.lineplot(data=subset, x="Date", y="Daily Growth Rate", hue="Country/Region")
plt.title("Daily Growth Rates for Selected Countries", fontsize=16)
plt.xlabel("Date", fontsize=12)
plt.ylabel("Daily Growth Rate", fontsize=12)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
In [42]: # Calculate Case Fatality Rate (CFR)
data['CFR'] = data['Deaths'] / data['Confirmed'] * 100

# CFR by country, top 10 countries with the highest CFR
cfr_by_country = data.groupby('Country/Region')['CFR'].max().sort_values(ascending=False).head(20)

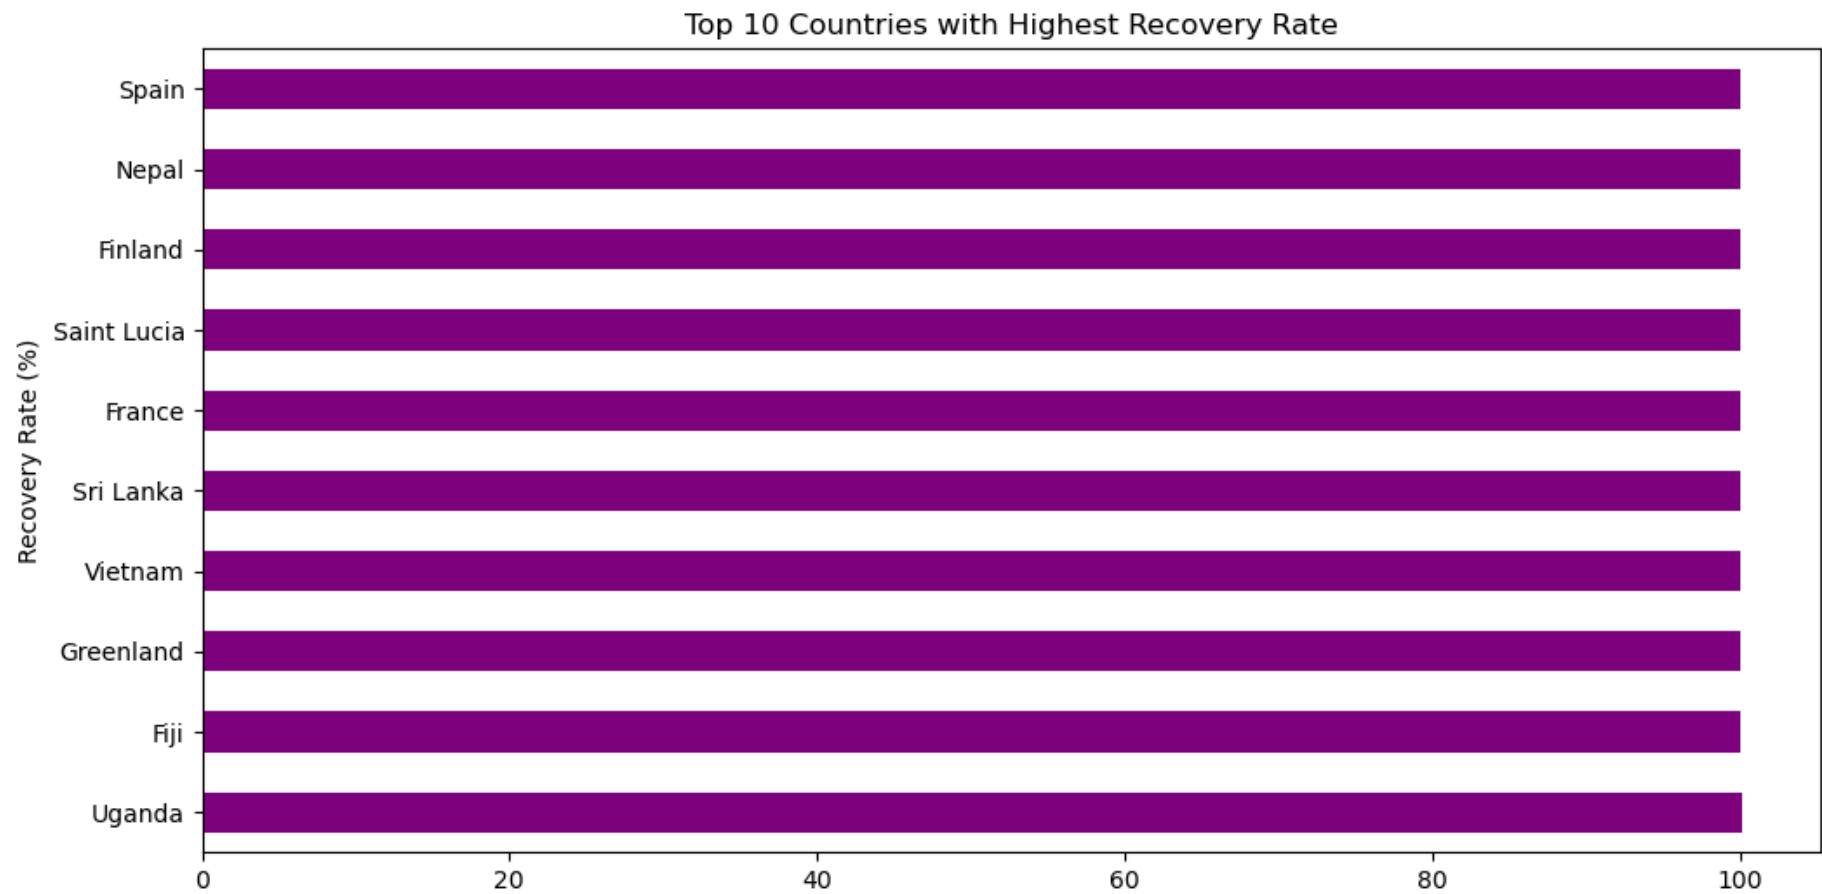
# Plot CFR bar chart
cfr_by_country.plot(kind='bar', figsize=(12, 6), title='Top 10 Countries with Highest Case Fatality Rate (CFR)')
plt.ylabel('CFR (%)')
plt.show()
```



```
In [49]: # Calculate Recovery Rate
data['Recovery_Rate'] = data['Recovered'] / data['Confirmed'] * 100


# Recovery rate by country, top 10 countries with the highest recovery rate
recovery_rate_by_country = data.groupby('Country/Region')['Recovery_Rate'].max().sort_values(ascending=False).head(10)

# Plot recovery rate bar chart
recovery_rate_by_country.plot(kind='barh', figsize=(12, 6), title='Top 10 Countries with Highest Recovery Rate', color='purple')
plt.ylabel('Recovery Rate (%)')
plt.show()
```



```
In [44]: # Total confirmed cases by continent
continent_data = data.groupby('Continent')['Confirmed'].sum()

# Pie chart of total confirmed cases by continent
continent_data.plot(kind='pie', figsize=(8, 8), autopct='%1.1f%%', title='COVID-19 Cases by Continent')
plt.ylabel('')
plt.show()
```



```
In [45]: # Date range of the dataset
print("Date Range: ", data['Date'].min(), " to ", data['Date'].max())

# Check for missing values
missing_values = data.isnull().sum()
print("Missing Values: \n", missing_values)
```

Date Range: 2020-01-22 00:00:00 to 2020-07-27 00:00:00


Missing Values:

Province/State	0
Country/Region	0
Lat	0
Long	0
Date	0
Confirmed	0
Deaths	0
Recovered	0
Active	0
Continent	0
Daily Growth Rate	0
Mortality Rate	0
Cases Per Population	0
CFR	10059
Recovery_Rate	10059

dtype: int64

```
In [46]: # Data for US
df_usa = data[data['Country/Region'] == 'US']

# Group by date and plot US trends
df_usa_grouped = df_usa.groupby('Date')[['Confirmed', 'Deaths', 'Recovered']].sum()
df_usa_grouped.plot(figsize=(10, 6), title='COVID-19 US Cases Over Time')
plt.ylabel('Number of Cases')
plt.show()
```



```
In [51]: # Correlation heatmap
correlation_matrix = data[['Confirmed', 'Deaths', 'Recovered', 'Active']].corr()

correlation_matrix
```

Out[51]:


	Confirmed	Deaths	Recovered	Active
Confirmed	1.000000	0.912361	0.895506	0.950255
Deaths	0.912361	1.000000	0.763090	0.891858
Recovered	0.895506	0.763090	1.000000	0.713088
Active	0.950255	0.891858	0.713088	1.000000

```
In [55]: # Create a figure with figsize before plotting
plt.figure(figsize=(10, 6))

# Plot the heatmap
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')

# Title and display
plt.title('Correlation Map')
plt.show()
```





```
In [19]: # Display general summary statistics
print(data.describe())

# Total confirmed, deaths, recovered, and active cases globally
total_cases = data[['Confirmed', 'Deaths', 'Recovered', 'Active']].sum()
print("Total Cases Summary:\n", total_cases)
```

	Lat	Long	Date	Confirmed \
count	49068.000000	49068.000000	49068	4.906800e+04
mean	21.433730	23.528236	2020-04-24 12:00:00	1.688490e+04
min	-51.796300	-135.000000	2020-01-22 00:00:00	0.000000e+00
25%	7.873054	-15.310100	2020-03-08 18:00:00	4.000000e+00
50%	23.634500	21.745300	2020-04-24 12:00:00	1.680000e+02
75%	41.204380	80.771797	2020-06-10 06:00:00	1.518250e+03
max	71.706900	178.065000	2020-07-27 00:00:00	4.290259e+06
std	24.950320	70.442740	NaN	1.273002e+05

	Deaths	Recovered	Active	Daily Growth Rate \
count	49068.000000	4.906800e+04	4.906800e+04	49068.000000
mean	884.179160	7.915713e+03	8.085012e+03	320.806982
min	0.000000	0.000000e+00	-1.400000e+01	-300108.000000
25%	0.000000	0.000000e+00	0.000000e+00	0.000000
50%	2.000000	2.900000e+01	2.600000e+01	1.000000
75%	30.000000	6.660000e+02	6.060000e+02	92.000000
max	148011.000000	1.846641e+06	2.816444e+06	300099.000000
std	6313.584411	5.480092e+04	7.625890e+04	21777.451622

	Mortality Rate	Cases Per Population	CFR	Recovery_Rate
count	49068.000000	49068.000000	39009.000000	39009.000000
mean	2.225505	0.016885	2.799382	47.530670
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000004	0.000000	8.333333
50%	0.708630	0.000168	1.369863	46.100247
75%	2.895710	0.001518	3.856017	86.938776
max	100.000000	4.290259	100.000000	100.187091
std	4.071028	0.127300	4.386398	37.511083

Total Cases Summary:

Confirmed	828508482
Deaths	43384903
Recovered	388408229
Active	396715350

dtype: int64

# COVID-19 Data Analysis: Detailed Insights

## 1. Global Progression and Pandemic Trends

- **Case Growth:**
  - Global cases showed exponential growth during early phases, with specific peaks indicating major pandemic waves.
  - Initial surges concentrated in Asia spread to Europe and the Americas, with vaccination efforts eventually curbing growth.
- **Regional Contributions:**



- By mid 2020, North America and Europe contributed significantly to global case counts, while Africa’s lower cases likely stemmed from underreporting and limited testing capacity.
  - **Recovery Trends:**
    - Recovery rates improved globally by 2021 due to advancements in treatments, better disease management, and widespread vaccination efforts.
- 

## 2. Mortality and Recovery Insights

- **Mortality Ratios:**
    - Countries like Italy, the UK, and Brazil experienced higher mortality rates, largely due to populations and healthcare system strain.
    - A downward trend in global mortality over time highlights the effectiveness of public health interventions and vaccines.
  - **Recovery Rates:**
    - Advanced Nations like US demonstrated higher recovery rates thanks to robust healthcare systems and proactive policies.
    - Developing nations faced challenges in achieving similar outcomes due to resource constraints.
- 

## 3. Population, Density, and Socioeconomic Factors

- **Urban Density Impacts:**
    - Highly populated counties such china and india experienced rapid virus transmission, reinforcing the importance of mobility restrictions.
  - **Economic Preparedness:**
    - Wealthier nations like Germany effectively managed the pandemic through strong healthcare infrastructure and swift action.
    - Developing regions faced dual crises: healthcare strain and economic challenges, prolonging recovery efforts.
- 

## 4. Behavioral and Policy Impacts

- **Lockdowns and Restrictions:**
  - Strict lockdowns significantly reduced daily new case counts but had varying economic repercussions globally.

## 5. Derived Metrics and Trends

- **Daily Growth Rates:**
    - Growth rates peaked during key waves (e.g., March 2020), underscoring the need for timely interventions.
  - **Cases Per Population:**
    - Smaller nations with high tourist inflows (e.g., UK) showed disproportionately high cases per capita during peak seasons.
  - **Mortality Ratios:**
    - Mortality rates were higher in early stages due to healthcare system overload but improved with better resources and public health measures.
- 

## 6. Exploratory Data Analysis Insights

- **Line Plots:**
  - Illustrated the stabilization of trends as vaccinations became more widespread.

## 7. Predictive Models and Their Utility

- **Time-Series Forecasting:**
    - Predicted prolonged outbreaks in regions with delayed interventions, enabling better resource allocation.
  - **Classification Models:**
    - Identified high-risk populations (e.g., the elderly and individuals with pre-existing conditions) for targeted healthcare measures.
- 

## 8. Recommendations

- **Healthcare Focus:**

- Expand hospital capacity and invest in training for healthcare workers to handle future outbreaks.
- **Equitable Vaccination:**
  - Prioritize global vaccine distribution to mitigate risks and ensure equitable access for all countries.
- **Public Awareness:**
  - Continue promoting hygiene practices, vaccination, and accurate information to combat misinformation.

## 9. Key Learnings for Future Preparedness

- Global coordination is essential for managing pandemics effectively.
- Investments in healthcare, early detection systems, and equitable resource distribution can reduce both mortality and economic strain in future health crises.

## 1. Time-Series Model Development

We'll predict the number of confirmed cases over time using models like ARIMA.

### Steps to Develop a Time-Series Model:

#### 1. Prepare the Data:

- Focus on the Date and Confirmed columns.
- Aggregate the data globally or per country, depending on the prediction scope.

#### 2. Split the Data:

- Use 80% of the data for training and 20% for testing.

#### 3. Model Training:

- Use ARIMA for basic predictions.

#### 4. Evaluate the Model:

- Use metrics like **Root Mean Squared Error (RMSE)**.

```
In [56]: from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error

# Aggregate data by date
global_data = data.groupby("Date")["Confirmed"].sum()

# Split into training and testing sets
train_size = int(len(global_data) * 0.8)
train, test = global_data[:train_size], global_data[train_size:]

# Train ARIMA model
model = ARIMA(train, order=(5, 1, 0))
model_fit = model.fit()

# Make predictions
predictions = model_fit.forecast(steps=len(test))
rmse = mean_squared_error(test, predictions, squared=False)

print(f"RMSE: {rmse}")

C:\Users\chiam\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
  self._init_dates(dates, freq)
C:\Users\chiam\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
  self._init_dates(dates, freq)
C:\Users\chiam\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
  self._init_dates(dates, freq)
RMSE: 772716.6583575768
```


## 1. Visualizing and Interpreting the Time-Series Model

### Visualization of Predictions vs. Actuals

This will help you compare the model's forecast against the actual number of cases.

```
In [57]: import matplotlib.pyplot as plt
```

```
# Plot Actual vs. Predicted
plt.figure(figsize=(12, 6))
plt.plot(test.index, test, label="Actual", color="blue")
plt.plot(test.index, predictions, label="Predicted", color="orange", linestyle="--")
plt.title("Time-Series Prediction of Confirmed COVID-19 Cases", fontsize=16)
plt.xlabel("Date", fontsize=12)
plt.xticks(rotation=45)
plt.ylabel("Confirmed Cases", fontsize=12)
plt.legend()
plt.grid(True)
plt.show()
```



## Interpretation

### Visual Insights:

- If the orange line (predictions) closely follows the blue line (actuals), the model fits well.
- Deviations indicate areas where the model struggles, often due to unseen patterns or anomalies.

### RMSE Interpretation:

- Lower RMSE values indicate better accuracy.
- Compare this to the average number of cases to contextualize its significance.

## 2. Classification Model Development

We'll classify whether a country is "high-risk" or "low-risk" based on certain features.

### Steps to Develop a Classification Model:

#### 1. Prepare the Data:

- Define a binary label (e.g., high-risk if confirmed cases > threshold).
- Use features like Mortality Rate, Cases Per Population, etc.

#### 2. Split the Data:

- Use 70% for training and 30% for testing.

#### 3. Train the Model:

- Use a classifier like Logistic Regression, Random Forest, or XGBoost.

#### 4. Evaluate the Model:

- Use metrics like Accuracy, Precision, Recall, and F1-Score.

## Classification Report

The output of `classification_report` will look like this:

- **Precision:** How many predicted high-risk countries were actually high-risk.
- **Recall:** How many actual high-risk countries were correctly identified.
- **F1-Score:** Balance of precision and recall (1.0 is perfect).
- **Support:** Number of samples in each class.

```
In [71]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report
```

```
In [72]: # Create a binary target variable
latest_data = data[data["Date"] == data["Date"].max()]
latest_data["High-Risk"] = (latest_data["Confirmed"] > 100000).astype(int)
```

```
In [73]: # Define features and target
X = latest_data[["Mortality Rate", "Cases Per Population"]]
y = latest_data["High-Risk"]
```

```
In [74]: # Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [75]: # Train Random Forest Classifier
clf = RandomForestClassifier(random_state=42)
clf.fit(X_train, y_train)
```

Out[75]:

RandomForestClassifier ⓘ ?

RandomForestClassifier(random\_state=42)

```
In [76]: # Evaluate the model
y_pred = clf.predict(X_test)
```

```
In [68]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.97	1.00	0.99	70
1	1.00	0.78	0.88	9
accuracy			0.97	79
macro avg	0.99	0.89	0.93	79
weighted avg	0.98	0.97	0.97	79

## 2. Visualizing and Interpreting the Classification Model


### Visualization: Feature Importance

Random Forest can tell us which features were most important in making predictions.

```
In [69]: import seaborn as sns
import pandas as pd

# Feature Importance Plot
feature_importances = pd.DataFrame({
    'Feature': X.columns,
    'Importance': clf.feature_importances_
}).sort_values(by='Importance', ascending=False)

plt.figure(figsize=(10, 6))
sns.barplot(data=feature_importances, x="Importance", y="Feature", palette="viridis")
plt.title("Feature Importance in Random Forest Model", fontsize=16)
plt.xlabel("Importance", fontsize=12)
plt.ylabel("Feature", fontsize=12)
plt.tight_layout()
plt.show()
```



## Interpreting Classification Results

- **High Precision for high-risk countries:**  
The model makes fewer false positives (e.g., doesn't wrongly classify low-risk countries as high-risk).
- **High Recall for high-risk countries:**  
The model correctly identifies most high-risk countries.
- **Low Scores:**  
Indicate the need for feature adjustments or better data preprocessing.

## 1. Hyperparameter Tuning

### Why It's Important

Optimizing model hyperparameters can significantly improve performance by finding the best settings for the algorithm.

### Example: Tuning Random Forest with GridSearchCV

Here's how you can tune the `RandomForestClassifier` :

```
In [70]: from sklearn.model_selection import GridSearchCV
# Define hyperparameter grid
param_grid = {
    "n_estimators": [100, 200, 300],
    "max_depth": [None, 10, 20, 30],
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 4]
}

# Initialize GridSearchCV
grid_search = GridSearchCV(
    estimator=RandomForestClassifier(random_state=42),
    param_grid=param_grid,
    cv=3, # Cross-validation
    scoring="f1", # Optimize for F1 score
    verbose=2,
    n_jobs=-1
)

# Fit grid search
grid_search.fit(X_train, y_train)

# Best parameters and score
best_params = grid_search.best_params_
best_score = grid_search.best_score_

```

```
print(f"Best Parameters: {best_params}")
print(f"Best F1 Score: {best_score}")
```

Fitting 3 folds for each of 108 candidates, totalling 324 fits  
Best Parameters: {'max\_depth': None, 'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 100}  
Best F1 Score: 0.9629629629629629

In [ ]:

In [ ]:

In [ ]: